

2023 554 R Notes on Spatial Data Analysis

Jon Wakefield
Departments of Biostatistics and Statistics
University of Washington

2023-01-06

R for Spatial Analysis

R has extensive spatial capabilities, the Spatial task view is here

Some of the notes that follow are build on Roger Bivand's notes taken from the latter site, and these are based on Bivand et al. (2013), which is a good reference book of developments until that point.

To get R code alone then load the `knitr` library and then type

```
purl("2023-554-Spatial-Classes.Rmd")
```

from the directory with this file in.

Overview of Spatial Classes

Class definitions are objects that contain the formal definition of a class of R objects, and are usually referred to as an S4 class.

Spatial classes were defined to represent and handle spatial data, so that data can be exchanged between different classes - they are different from regular classes since they need to contain information about spatial locations and their coordinate reference systems

- The `sp` library was traditionally the workhorse for representing spatial data.
- Now `sf` is being increasingly used, details are here

These notes will focus on `sf`.

See Spatial Data Science with Applications in R by Pebesma and Bivand, for a more up to date version.

Another resource is the book Geocomputation with R

The `sf` Representation

The `sf` package builds on R packages used in spatial data including `sp`, `rgeos` and the vector parts of `rgdal`, providing an interface to certain tidyverse packages.

This package reads and writes data through GDAL, and uses GEOS, s2geometry, and PROJ.

The `sf` package has `sf` objects, a sub-class of a `data.frame` or a `tibble`. These objects contain at least one geometry list-column of class `sfc`, where each element contains the geometry as an R object of class `sfg`.

All functions are prefixed with `st_` to indicate “spatial type” to make them more easily searchable on the command line.

You can load the package by running `library(sf)`, but note that if you’ve never used this package before, you should run `install.packages('sf')` in your R console, and then load the package with the `library(sf)` command.

Creating a Spatial Object

We can use the `st_read` function (part of the `sf` package) to read in a spatial file.

```
file <- system.file("gpkg/nc.gpkg", package = "sf")
nc <- st_read(file)
## Reading layer `nc.gpkg' from data source
##   `/Library/Frameworks/R.framework/Versions/4.2/Resources/library/sf/gpkg/nc.gpkg'
##   using driver `GPKG'
## Simple feature collection with 100 features and 14 fields
## Geometry type: MULTIPOLYGON
## Dimension:      XY
## Bounding box:   xmin: -84.32385 ymin: 33.88199 xmax: -75.45698 ymax: 36.58965
## Geodetic CRS:   NAD27
```

Note that the `st_read()` function takes in two arguments, the `dsn`, or data source name, and the layer. Our provided data has only one layer, so that argument is omitted in our example. You can check the available layers by querying:

```
st_layers(file)
## Driver: GPKG
## Available layers:
##   layer_name geometry_type features fields
## 1 nc.gpkg Multi Polygon      100      14
```

It’s also possible to subset this data and plot it. You can run each line sequentially in R to see what each line of code adds to the plot. Note that this is written in base R, not with `ggplot()` where each line of code builds upon the plot.

```
# Subset the data
nc5 <- nc[1:5, ]
nc7 <- nc[1:7, ]

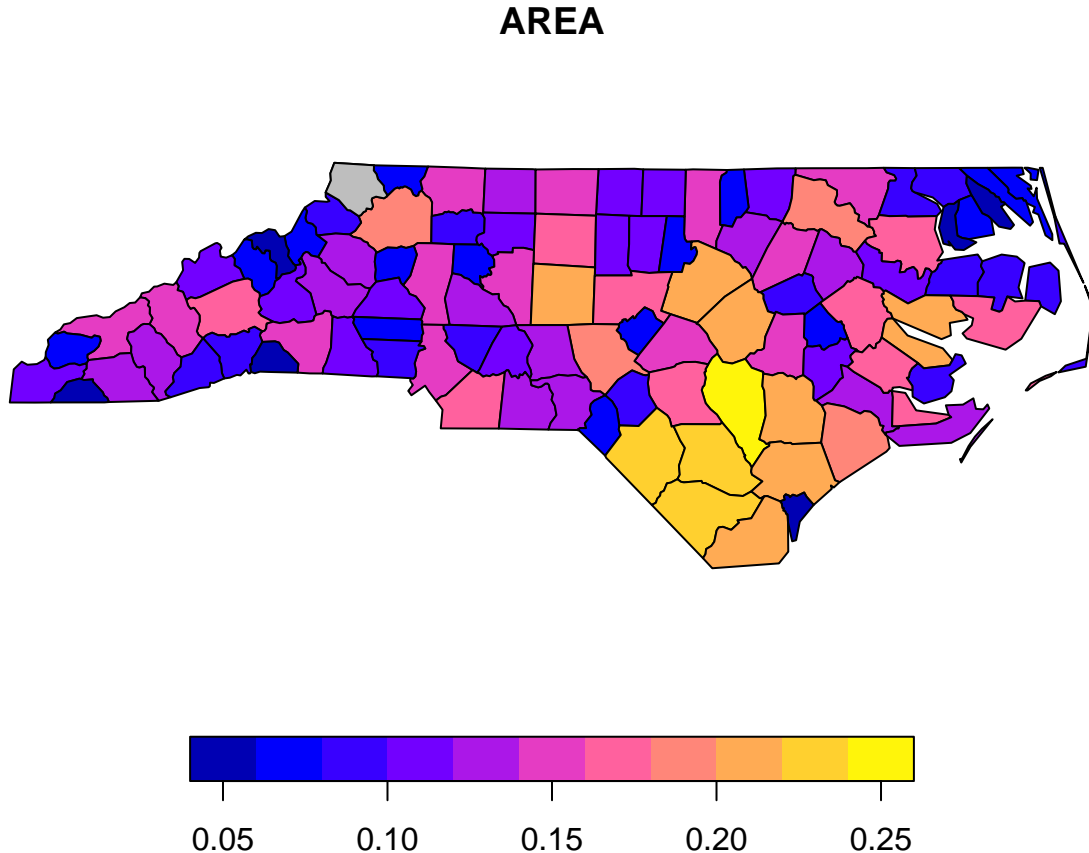
# Plot the data
plot(st_geometry(nc7))
plot(st_geometry(nc5), add = TRUE, border = "brown")
cc <- st_coordinates(st_centroid(st_geometry(nc7)))
text(cc, labels = 1:nrow(nc7), col = "blue")
```



Visualizing Spatial Data


As you can see below, we can create a plot of the sizes in area of each county in North Carolina. The area part comes into play when we index the `nc` object with a 1. If you change this to a 2, you'll see that this now plots the perimeter of each county instead, as perimeter is the second feature.

```
par(mar = c(0,0,1,0))  
plot(nc[1], reset = FALSE) # reset = FALSE: we want to add to a plot with a legend  
plot(nc[1,1], col = 'grey', add = TRUE)
```



We can also look at the `meuse` dataset from the `sp` package and convert the data into a `sp` object, namely a `SpatialPointsDataFrame` by using the `coordinates()` function. Then, using `st_as_sf()`, we can go from an `sp` object to an `sf` object.

```
library(sp)  
data(meuse)  
coordinates(meuse) = ~x+y  
m.sf = st_as_sf(meuse)  
opar = par(mar=rep(0,4))  
plot(m.sf)
```



John Snow Example


For fun, let's look at the poster child of health mapping.

The Snow data consists of the relevant 1854 London streets, the location of 578 deaths from cholera, and the position of 13 water pumps (wells) that can be used to re-create John Snow's map showing deaths from cholera in the area surrounding Broad Street, London in the 1854 outbreak.

```
library(HistData)
data(Snow.deaths)
data(Snow.pumps)
data(Snow.streets)
# data(Snow.polygons)
```

We first create a `SpatialLines` object containing the coordinates of the streets using the `Lines` function


```
# Streets
slist <- split(Snow.streets[,c("x", "y")],
              as.factor(Snow.streets[, "street"]))
Ll1 <- lapply(slist, Line)
Lsl1 <- Lines(Ll1, "Street")
Snow.streets.sp <- SpatialLines(list(Lsl1))
m <- st_as_sf(Snow.streets.sp)
plot(m)
```



Display the streets and then add the deaths and pumps (with labels).

The red squares are deaths, blue triangles are pumps

```
plot(m, col="gray")
# deaths
Snow.deaths.sp = SpatialPoints(Snow.deaths[,c("x", "y")])
p <- st_as_sf(Snow.deaths.sp)
plot(p, add=TRUE, col='red', pch=15, cex=0.6)
# pumps
spp <- SpatialPoints(Snow.pumps[,c("x", "y")])
Snow.pumps.sp <-
  SpatialPointsDataFrame(spp, Snow.pumps[,c("x", "y")])
spp <- st_as_sf(Snow.pumps.sp)
plot(spp, add=TRUE, col='blue', pch=17, cex=1.0)
text(Snow.pumps[,c("x", "y")], labels=Snow.pumps$label,
     pos=1, cex=0.8)
```



Reading Shapefiles

ESRI (a company one of whose products is ArcGIS) shapefiles consist of three files, and this is a common form.

- The first file (*.shp) contains the geography of each shape.
- The second file (*.shx) is an index file which contains record offsets.
- The third file (*.dbf) contains feature attributes with one record per feature.

The Washington state Geospatial Data Archive can be accessed here contains data that can be read into R.

As an example, consider ashington county data that was downloaded from wagda.

The data consists of the three files: wacounty.shp, wacounty.shx, wacounty.dbf.

The following code reads in these data and then draws a county level map of 1990 populations, and a map with centroids.

First load the libraries.

```
library(maps)
library(shapefiles)
library(maptools)
```

Note that there are problems with the files, which are sorted by using the `repair=T` argument.

The data can be downloaded from here: [here](#)

```
library(rgdal) # for readOGR
library(spdep) # for poly2nb and nb2inla
download.file("http://faculty.washington.edu/jonno/SISMIDmaterial/wacounty.shp",destfile = "wacounty.shp")
download.file("http://faculty.washington.edu/jonno/SISMIDmaterial/wacounty.shx",destfile = "wacounty.shx")
download.file("http://faculty.washington.edu/jonno/SISMIDmaterial/wacounty.dbf",destfile = "wacounty.dbf")
wacounty=readOGR(dsn=".",layer = "wacounty")
## OGR data source with driver: ESRI Shapefile
## Source: "/Users/andreaboskovic/Desktop/uw classes/year2/winter 23/554/dropbox materials/updated R no
## with 39 features
## It has 6 fields
class(wacounty)
## [1] "SpatialPolygonsDataFrame"
## attr(,"package")
## [1] "sp"
```

Let's see what these variables look like: we see county names and FIPS codes.

```
names(wacounty)
## [1] "AreaName" "AreaKey" "INTPTLAT" "INTPTLNG" "TotPop90" "CNTY"
wacounty$AreaName[1:3]
## [1] "WA, Adams County" "WA, Asotin County" "WA, Benton County"
wacounty$AreaKey[1:3]
## [1] "53001" "53003" "53005"
head(wacounty)
## class      : SpatialPolygonsDataFrame
```

```
## features      : 6
## extent       : -124.7312, -116.915, 45.5434, 48.5504 (xmin, xmax, ymin, ymax)
## crs          : NA
## variables     : 6
## names        :      AreaName, AreaKey,  INTPTLAT,   INTPTLNG, TotPop90, CNTY
## min values   : WA, Adams County,   53001, 45.773673, -123.931203,   13603,   1
## max values   : WA, Clark County,   53011, 48.109502, -117.18502,   238053,   9
```

```
head(wacounty@data)
##      AreaName AreaKey INTPTLAT  INTPTLNG TotPop90 CNTY
## 0  WA, Adams County  53001 46.98899 -118.5569   13603   1
## 1  WA, Asotin County  53003 46.18248 -117.1850   17605   3
## 2  WA, Benton County  53005 46.24764 -119.5015  112560   5
## 3  WA, Chelan County  53007 47.87696 -120.6414   52250   7
## 4  WA, Clallam County  53009 48.10950 -123.9312   56464   9
## 5  WA, Clark County  53011 45.77367 -122.4843  238053  11
```


Drawing a map

We look at some variables.

```
wacounty$INTPTLAT[1:3] # latitude
## [1] 46.98899 46.18248 46.24764
wacounty$INTPTLNG[1:3] # longitude
## [1] -118.5569 -117.1850 -119.5015
wacounty$CNTY[1:3]
## [1] "1" "3" "5"
wacounty$TotPop90[1:3]
## [1] 13603 17605 112560
```

We look at some variables, and then set up the colors to map. We map 1990 Washington population counts by census tracts.


```
plotvar <- wacounty$TotPop90 # variable we want to map
summary(plotvar)
##      Min. 1st Qu.  Median    Mean 3rd Qu.    Max.
##      2248  16863   38558 123102   81478 1507319
nclr <- 8 # next few lines set up the color scheme for plotting
plotclr <- brewer.pal(nclr,"Oranges")
brks <- round(quantile(plotvar,probs=seq(0,1,1/(nclr))),digits=1)
colnum <- findInterval(plotvar,brks,all.inside=T)
colcode <- plotclr[colnum]
plot(wacounty,col=colcode)
legend(-119,46,legend=leglabs(round(brks,digits=1)),fill=plotclr,cex=0.35,bty="n")
```



In ggplot:


```
wacounty_fort <- fortify(wacounty, region = "AreaKey")
# cbind in plotvar and colcode
temp_df <- data.frame(id = wacounty$AreaKey, plotvar = wacounty$TotPop90, colcode = colcode)
wacounty_fort <- left_join(wacounty_fort, temp_df)

ggplot() +
  geom_polygon(data = wacounty_fort, aes(x = long, y = lat, group = id, fill = colcode), col = "black")
  scale_fill_manual(values = rev(plotclr), labels = rev(leglabs(round(brks,digits=1)))) +
  theme_bw() +
  coord_fixed() +
  xlab("") +
  ylab("") +
  theme(legend.title = element_blank(),
        panel.grid = element_blank(),
        panel.border = element_blank(),
        axis.ticks = element_blank(),
        axis.text = element_blank())
```




As an alternative we can use the `spplot` function, which uses lattice (trellis) plot methods for spatial data with attributes. We map the 1990 Washington population counts by county.


```
spplot(wacounty, zcol="TotPop90")
```




In ggplot:

```
ggplot() +
  geom_polygon(data = wacounty_fort, aes(x = long, y = lat, group = id, fill = plotvar), col = "black") +
  theme(legend.title = element_blank(),
        panel.grid = element_blank(),
        panel.border = element_blank(),
        axis.ticks = element_blank(),
        axis.text = element_blank())
```



For illustration, we define our own cutpoints by hand.


```
summary(wacounty$TotPop90)
##      Min. 1st Qu.  Median    Mean 3rd Qu.    Max.
##      2248  16863   38558 123102   81478 1507319
wacounty$cats <- cut(wacounty$TotPop90, breaks=c(0,17000,39000,100000,1520000), labels=c("<17,000", "17,000", "39,000", "100,000", "1,520,000"))
spplot(wacounty, zcol="cats",
       col.regions=brewer.pal(4, "Blues"))
```



In ggplot:

```
# create variable for which bin a county falls in
wacounty_fort$col_discrete <- factor(ifelse(wacounty_fort$plotvar < 17000, "<17,000",
  ifelse(wacounty_fort$plotvar < 39000, "17,000-39,000",
    ifelse(wacounty_fort$plotvar < 100000, "39,000-100,000", ">100,000"))), levels=c("<17,000", "17,000-39,000", "39,000-100,000", ">100,000"))

ggplot() +
  geom_polygon(data = wacounty_fort, aes(x = long, y = lat, group = id,
    fill = col_discrete), col = "black") +
  theme_bw() +
  coord_fixed() +
  xlab("") +
  ylab("") +
  theme(legend.title = element_blank(),
    panel.grid = element_blank(),
    panel.border = element_blank(),
    axis.ticks = element_blank(),
    axis.text = element_blank()) +
  scale_fill_manual(values = (brewer.pal(4, "Blues")))
```



We now highlight a county


```
# identify counties of interest
xx=which(wacounty$CNTY==33)
xx2=which(wacounty$CNTY==63)
# plot the whole state
plot(wacounty, border="#0000075")
```

```

# highlight counties of interest
plot(wacounty[xx, ], col="#ff000075", add=T)
plot(wacounty[xx2, ], col=NA, border="#0000ff75", add=T, lwd=2.5)

# Add some labels
text(coordinates(wacounty[xx, ]), "King", cex=0.75, pos=3, offset=0.25)
text(coordinates(wacounty[xx2, ]), "Spokane", cex=0.7, pos=1, offset=0.25)
points(coordinates(wacounty[c(xx, xx2), ]), pch=16, cex=0.75)

```



Now let's repeat for census tracts.

```

## Reading Shapefiles
library(rgdal) # for readOGR
library(spdep) # for poly2nb and nb2inla
download.file("http://faculty.washington.edu/jonno/SISMIDmaterial/watract.shp", destfile = "wattract.shp")
download.file("http://faculty.washington.edu/jonno/SISMIDmaterial/watract.shx", destfile = "wattract.shx")
download.file("http://faculty.washington.edu/jonno/SISMIDmaterial/watract.dbf", destfile = "wattract.dbf")
wattract=readOGR(dsn=".", layer = "wattract")
## OGR data source with driver: ESRI Shapefile
## Source: "/Users/andreaboskovic/Desktop/uw classes/year2/winter 23/554/dropbox materials/updated R no
## with 1152 features
## It has 7 fields

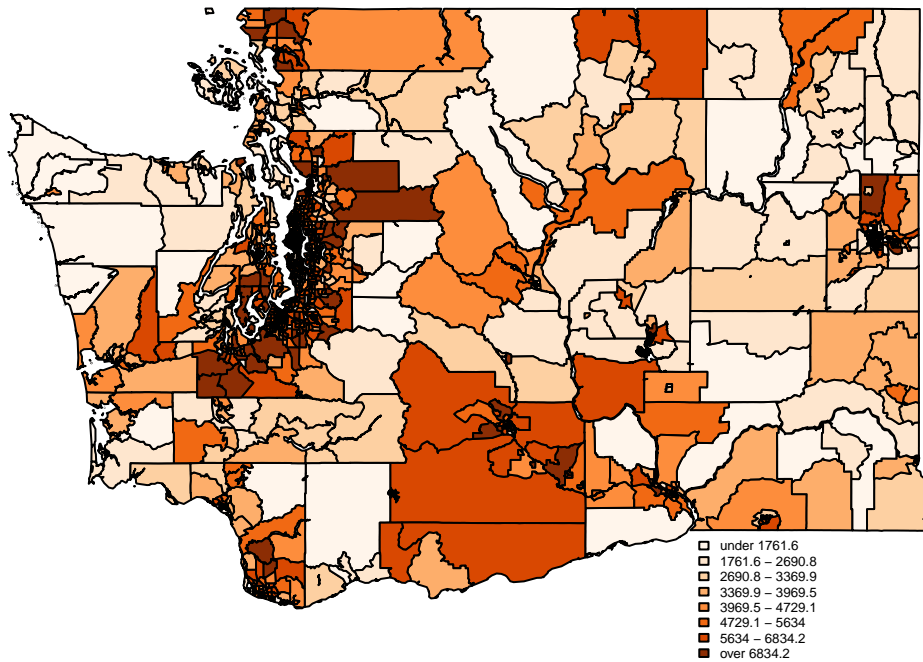
```

Drawing a census tract map: we repeat but now map populations at the census tract level.

```

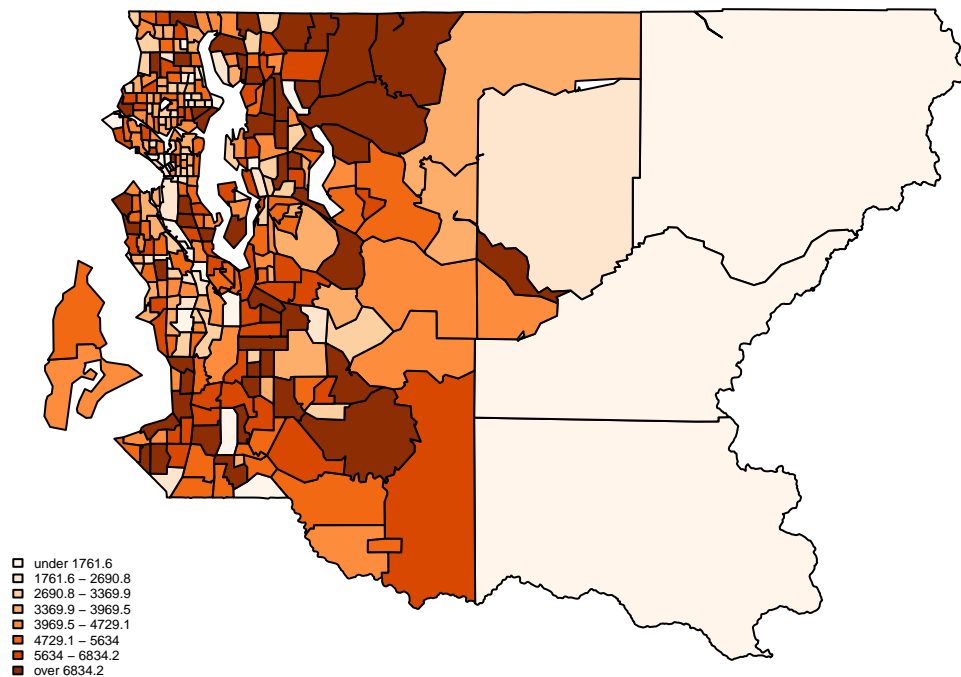
wattract <- readShapePoly(fn="wattract",proj4string=
  CRS("+proj=longlat"),repair=T) #
names(wattract)
## [1] "AreaName" "AreaKey" "INTPTLAT" "INTPTLNG" "TotPop90" "TRACT" "CNTY"
plotvar <- wattract$TotPop90 # variable we want to map
brks <- round(quantile(plotvar,probs=seq(0,1,1/(nclr))),digits=1)
colnum <- findInterval(plotvar,brks,all.inside=T)
colcode <- plotclr[colnum]
plot(wattract,col=colcode)
legend(-119,46,legend=leglabs(round(brks,digits=1)),
      fill=plotclr,cex=0.4,bty="n")

```



We zoom in on King County

```
xx=which(watract$CNTY==33)
plot(watract[xx,], col=colcode[xx])
legend("bottomleft", legend=leglabs(round(brks,digits=1)),
      fill=plotclr,cex=0.4,bty="n")
```

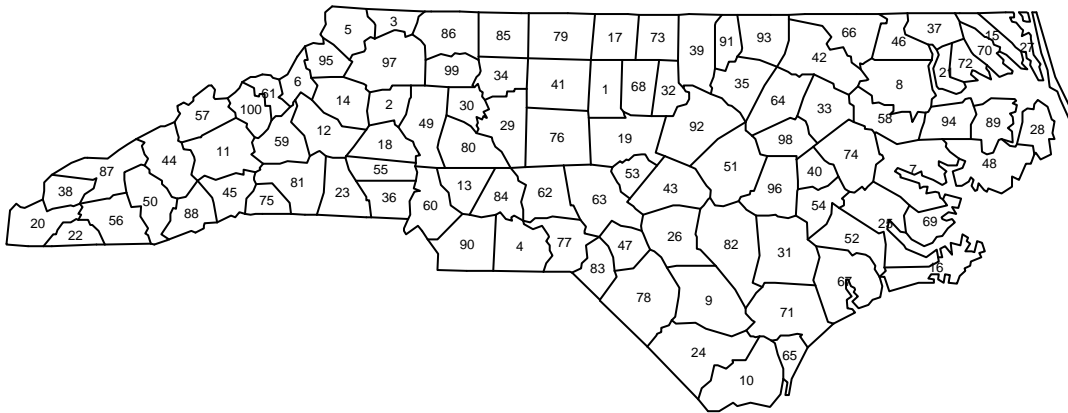


A county map of North Carolina with text

```

library(ggplot2) # For map_data. It's just a wrapper; should just use maps.
library(sp)
library(maps)
getLabelPoint <- # Returns a county-named list of label points
function(county) {Polygon(county[c('long', 'lat')])@labpt}
df <- map_data('county', 'north carolina') # NC region county data
centroids <- by(df, df$subregion, getLabelPoint) # Returns list
centroids <- do.call("rbind.data.frame", centroids) # Convert to Data Frame
names(centroids) <- c('long', 'lat') # Appropriate Header
map('county', 'north carolina')
text(centroids$long, centroids$lat, rownames(centroids), offset=0, cex=0.4)

```

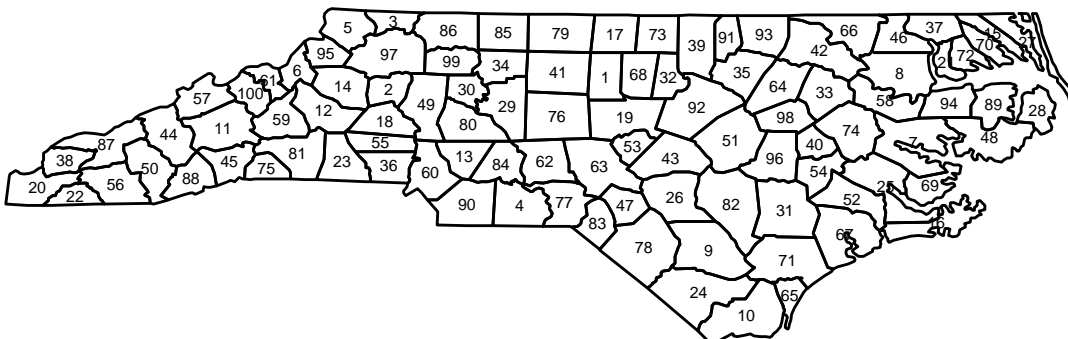


In ggplot:

```

ggplot() +
  geom_polygon(data = df, aes(x = long, y = lat, group = group), fill = "NA", col = "black") +
  geom_text(aes(x = centroids[,1], y = centroids[,2], label = 1:100), size = 2) +
  theme_bw() +
  coord_fixed() +
  xlab("") +
  ylab("") +
  theme(legend.title = element_blank(),
        panel.grid = element_blank(),
        panel.border = element_blank(),
        axis.ticks = element_blank(),
        axis.text = element_blank())

```



Below is a map using `sf`. It maps the variable `BIR74` which is the number of births in 1974, by counties in North Carolina.

```
library(tidyverse)
library(sf)
system.file("gpkg/nc.gpkg", package = "sf") %>%
  read_sf() -> nc
nc.32119 <- st_transform(nc, 'EPSG:32119')
nc.32119 %>%
  select(BIR74) %>%
  plot(graticule = TRUE, axes = TRUE)
```

